

Utilization alone is not an access summary.

Utilization measures completed visits per **slot**.

Served rate measures completed visits per **arrival**.

Minimum reporting set: served rate, utilization, offered wait, no-offer share, outcome decomposition, class-level served rates.

Under high demand, high utilization can coexist with poor access.

Stress-test baseline: $\lambda = 100$ arrivals/day, $S = 32$ slots/day.

Utilization ≈ 0.841 ; served rate ≈ 0.269 .

The baseline is useful for mechanism diagnosis, not as a calibrated clinic estimate.

Balking reallocates access; it does not recover aggregate capacity.

Higher balking reduces own-class served rate.

Through pooled FCFS, the other class can gain access.

Removing balking shifts losses toward no-offer and cancellation.

Cancellation can be reabsorbed; no-show capacity is lost.

Cancellation releases future slots that other patients may fill.

No-show slots are not rebooked.

No-show behavior is the most direct behavioral driver of utilization.

Aggregate metrics can hide class-level access gaps.

Symmetric parameters imply near-zero expected class gap.

Asymmetric behavior creates class gaps under pooled FCFS.

Meeting decision: stress-test baseline for mechanism diagnosis, or calibrated realistic scenario for operational interpretation.